

Dataset and Preprocessing


I am using the MNIST dataset to evaluate and compare model architecture and parameters. Both networks presented take an image (28x28 pixels) of a handwritten digit (between 0 and 9) and it will try to predict its class. Below is an example of a handwritten letter and the correct label.



After loading the MNIST dataset, all images were normalized from 0–255 pixel data to 0-1 to allow for faster learning of input node parameters (normalized to avoid exploding gradients). Before using the dataset, the labels (0-9) had to be converted from ordinal values to sets (this mean that value 1 is not smaller than 2, they are just two different sets of classification), this was done by one hot encoding the labels, resulting in 10 output nodes which was easier to code than how thescikit-learn library handles it.

ANN Design

- Both models have the same input nodes. Input layer is defined as a vector with 784 entries (from the 28x28 image).
- This model has 1 hidden layer with 64 nodes and uses the sigmoid activation function
- The output layer is 10 nodes corresponding to numbers, it layer uses the softmax activation function to normalize all values between 0 and 1 (sum of 1), this way output values exist as probabilities and the largest probability is chosen as the prediction.
- Backpropagation or gradient descent was used to find best results (minima of curve). By implementing this algorithm, possibly underfitted graphs could be fitted more optimally.



Model 1 – No Library Multilayer Neural Net

- **Initial weights** = Initial weights and biases were randomized. I chose to randomly initialize all weights since the data has been normalized and the spread is in a very small distribution. The initial weights can be a factor in accuracy since a model can converge to a local minima, but I feel that this is mitigated by normalizing the dataset and training until the model converges.
- **Hidden layer** = Since the data is not linearly separable, a hidden layer is required. I avoided adding more hidden layers as one is usually sufficient for majority of simple problems.
- **Nodes** = I initially started with only 12 neurons in the hidden layer, keeping all other parameters the same, it seemed the model was underfitting because the training accuracies were quite low. I then increased the number of nodes to 32 and then 64. In my experience 64 nodes was enough to observe a significant improvement in accuracy and sufficient enough as loss started to level off.
- **Learning rate** = This is a parameter of gradient descent, learning rate informs the model on the level of change depending on the estimated error after weight update. I first chose a very small learning rate of 0.001, however it resulted in a very long training time, therefore I chose the value of 1 as it is also not too big where it would result in sub-optimal set of weights.
- **Momentum** = I decided to add momentum which replaces gradient with a momentum which is an aggregation of all gradients, after adding momentum and trying various values I noticed a decent improvement in training time.
- **Epochs** = When a dataset, such as this one, is large, epochs and batches can be used to divide the data into smaller sizes and given to the model in steps and updates the weights at the end of every step to fit the given data, therefore only having one epoch could lead to underfitting. However, having too many epochs could also lead to overfitting so I used various values to observe changes in model performance. It seemed that after 200-250 epochs the model could be overfitting, so I set the epoch number to 200.
- **Batches** = Batch represents the subset of training examples in a single batch before the model's internal parameters are updated. I decided to include batch in my network because it requires less memory and improves gradient descent. A large batch number will make coarse updates to the weights whereas a smaller batch number implies the model updates often. I noticed that a very large batch size (1000) decreased training time significantly, however it also does much worse with test data compared to smaller batch numbers. Another factor I noticed was that using very high batch numbers took much longer for loss to converge. Using batch, allows the model to converge faster and avoid excessive number of epochs if correctly set.
- **Weight updates** = instead of loops, I used matrix multiplication during back propagation to reduce training time. Doing so reduced my model training time from approximately 90+ minutes to 1-5 minutes

Model 1 (100 epochs)

Final loss: 0.7298848203148632										Final loss: 0.6927345992527022									
Train data stats:										Test data stats:									
[[5131 1 179 89 40 267 126 80 116 54]										[[867 0 31 13 5 43 32 10 21 9]									
[1 6293 135 48 21 45 46 148 207 17]										[0 1052 31 2 3 4 3 27 11 5]									
[116 133 4376 223 99 111 287 170 230 65]										[21 19 763 30 13 24 52 31 25 11]									
[83 14 247 4592 28 538 18 72 480 124]										[7 4 53 806 2 95 2 19 93 14]									
[16 12 165 34 4480 155 140 221 169 491]										[5 1 26 4 762 23 33 25 29 82]									
[242 28 65 522 63 3445 173 85 399 141]										[40 3 10 78 7 591 27 10 57 27]									
[190 12 265 30 220 188 4969 36 134 62]										[25 2 43 2 35 27 786 3 28 6]									
[53 29 179 75 114 47 26 4976 77 490]										[5 5 28 12 17 7 6 815 16 58]									
[78 204 307 415 188 451 82 100 3715 167]										[9 48 41 52 27 59 12 28 627 13]									
[13 16 40 103 589 174 51 377 324 4338]]										[1 1 6 11 111 19 5 60 67 784]]									
precision					recall					precision					recall				
0	0.87	0.84	0.85	6083	0	0.88	0.84	0.86	1031	0	0.88	0.84	0.86	1031	0	0.88	0.84	0.86	1031
1	0.93	0.90	0.92	6961	1	0.93	0.92	0.93	1138	1	0.93	0.92	0.93	1138	1	0.93	0.92	0.93	1138
2	0.73	0.75	0.74	5810	2	0.74	0.77	0.76	989	2	0.74	0.77	0.76	989	2	0.74	0.77	0.76	989
3	0.75	0.74	0.75	6196	3	0.80	0.74	0.77	1095	3	0.80	0.74	0.77	1095	3	0.80	0.74	0.77	1095
4	0.77	0.76	0.76	5883	4	0.78	0.77	0.77	990	4	0.78	0.77	0.77	990	4	0.78	0.77	0.77	990
5	0.64	0.67	0.65	5163	5	0.66	0.70	0.68	850	5	0.66	0.70	0.68	850	5	0.66	0.70	0.68	850
6	0.84	0.81	0.83	6106	6	0.82	0.82	0.82	957	6	0.82	0.82	0.82	957	6	0.82	0.82	0.82	957
7	0.79	0.82	0.81	6066	7	0.79	0.84	0.82	969	7	0.79	0.84	0.82	969	7	0.79	0.84	0.82	969
8	0.63	0.65	0.64	5707	8	0.64	0.68	0.66	916	8	0.64	0.68	0.66	916	8	0.64	0.68	0.66	916
9	0.73	0.72	0.72	6025	9	0.78	0.74	0.76	1065	9	0.78	0.74	0.76	1065	9	0.78	0.74	0.76	1065
accuracy					accuracy					accuracy					accuracy				
macro avg					macro avg					macro avg					macro avg				
weighted avg					weighted avg					weighted avg					weighted avg				

(left is train, right is test)

Model 1 (200 epochs)

Final loss: 0.5518701951222055										Final loss: 0.5264352336787862									
Train data stats:										Test data stats:									
[[5343 1 124 69 28 197 100 56 69 42]										[[901 0 26 8 4 38 27 7 17 6]									
[1 6400 92 41 15 43 37 115 166 15]										[0 1073 16 3 2 1 3 27 7 4]									
[71 70 4777 201 64 82 175 147 164 62]										[14 8 841 29 9 14 24 34 21 7]									
[56 18 182 4919 18 416 13 49 355 112]										[3 4 33 851 1 69 1 12 71 9]									
[12 11 153 19 4855 108 107 142 109 399]										[2 2 22 2 826 21 26 18 19 70]									
[191 30 48 385 44 3933 133 44 318 99]										[32 2 4 61 7 662 25 4 54 22]									
[136 11 193 31 165 153 5257 16 90 32]										[18 3 31 4 21 25 841 2 23 5]									
[41 26 143 69 67 39 17 5318 59 309]										[3 3 23 9 10 5 2 853 18 42]									
[57 156 212 320 121 319 60 65 4284 144]										[7 39 30 35 17 46 6 21 700 16]									
[15 19 34 77 465 131 19 313 237 4735]]										[0 1 6 8 85 11 3 50 44 828]]									
precision					recall					precision					recall				
0	0.90	0.89	0.89	6029	0	0.92	0.87	0.89	1034	0	0.92	0.87	0.89	1034	0	0.92	0.87	0.89	1034
1	0.95	0.92	0.94	6925	1	0.95	0.94	0.94	1136	1	0.95	0.94	0.94	1136	1	0.95	0.94	0.94	1136
2	0.80	0.82	0.81	5813	2	0.81	0.84	0.83	1001	2	0.81	0.84	0.83	1001	2	0.81	0.84	0.83	1001
3	0.80	0.80	0.80	6138	3	0.84	0.81	0.82	1054	3	0.84	0.81	0.82	1054	3	0.84	0.81	0.82	1054
4	0.83	0.82	0.83	5915	4	0.84	0.82	0.83	1008	4	0.84	0.82	0.83	1008	4	0.84	0.82	0.83	1008
5	0.73	0.75	0.74	5225	5	0.74	0.76	0.75	873	5	0.74	0.76	0.75	873	5	0.74	0.76	0.75	873
6	0.89	0.86	0.88	6084	6	0.88	0.86	0.87	973	6	0.88	0.86	0.87	973	6	0.88	0.86	0.87	973
7	0.85	0.87	0.86	6088	7	0.83	0.88	0.85	968	7	0.83	0.88	0.85	968	7	0.83	0.88	0.85	968
8	0.73	0.75	0.74	5738	8	0.72	0.76	0.74	917	8	0.72	0.76	0.74	917	8	0.72	0.76	0.74	917
9	0.80	0.78	0.79	6045	9	0.82	0.80	0.81	1036	9	0.82	0.80	0.81	1036	9	0.82	0.80	0.81	1036
accuracy					accuracy					accuracy					accuracy				
macro avg					macro avg					macro avg					macro avg				
weighted avg					weighted avg					weighted avg					weighted avg				

Model 1 (200 batch, 20 epoch)

Final loss: 0.1666956036373846					Final loss: 0.20033133006317197				
Train data stats:					Test data stats:				
[[5797 0 27 9 7 32 30 8 16 21]					[[963 0 9 1 1 8 10 1 10 6]				
[0 6598 28 12 7 10 7 21 41 9]					[0 1115 2 4 2 1 3 11 3 6]				
[12 45 5648 82 26 23 28 48 38 12]					[1 3 965 13 8 2 4 19 7 2]				
[8 15 51 5754 6 106 2 17 86 53]					[2 2 16 949 1 34 1 14 18 8]				
[7 9 34 5 5570 16 32 27 16 134]					[0 0 7 1 914 4 7 5 9 26]				
[21 7 15 125 7 5062 53 7 81 25]					[4 1 1 21 0 792 8 1 17 8]				
[33 6 30 5 42 54 5736 2 34 4]					[8 5 11 1 16 18 923 0 14 1]				
[7 20 47 38 20 10 0 6015 12 98]					[2 2 6 7 3 4 0 947 4 9]				
[33 27 65 66 25 73 28 16 5478 40]					[0 7 13 10 4 21 2 3 880 5]				
[5 15 13 35 132 35 2 104 49 5553]]					[0 0 2 3 33 8 0 27 12 938]]				
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.98	0.97	0.98	5947	0	0.98	0.95	0.97	1009
1	0.98	0.98	0.98	6733	1	0.98	0.97	0.98	1147
2	0.95	0.95	0.95	5962	2	0.94	0.94	0.94	1024
3	0.94	0.94	0.94	6098	3	0.94	0.91	0.92	1045
4	0.95	0.95	0.95	5850	4	0.93	0.94	0.94	973
5	0.93	0.94	0.94	5403	5	0.89	0.93	0.91	853
6	0.97	0.96	0.97	5946	6	0.96	0.93	0.94	997
7	0.96	0.96	0.96	6267	7	0.92	0.96	0.94	984
8	0.94	0.94	0.94	5851	8	0.90	0.93	0.92	945
9	0.93	0.93	0.93	5943	9	0.93	0.92	0.92	1023
accuracy			0.95	60000	accuracy			0.94	10000
macro avg	0.95	0.95	0.95	60000	macro avg	0.94	0.94	0.94	10000
weighted avg	0.95	0.95	0.95	60000	weighted avg	0.94	0.94	0.94	10000

(left is train, right is test)

Model 2 – scikit-learn MLP Classifier

- First, I ran the MLP model with the same parameters (shown below) that I picked for the above model.

```
# network parameters
n_x = X_train.shape[0] # num input nodes
n_h = 64                # num hidden layer nodes
learning_rate = 1       # learning rate
epochs = 20             # epochs
momentum = 0.5          # momentum
```

- Epoch was set to 20 at first because MLP by default implements 200 batches just like model 1 (with batch). But I had to increase it to 50 because the solution wasn't optimal

Train data stats:					Test data stats:				
[[5923 0 0 0 0 0 0 0 0 0]					[[968 0 6 0 1 3 6 1 6 1]				
[0 6742 0 0 0 0 0 0 0 0]					[1 1122 3 0 0 0 2 6 0 2]				
[0 0 5958 0 0 0 0 0 0 0]					[1 4 1001 8 4 0 4 10 4 0]				
[0 0 0 6131 0 0 0 0 0 0]					[0 2 5 978 1 11 1 3 4 3]				
[0 0 0 0 5842 0 0 0 0 0]					[0 0 4 1 961 0 6 1 5 7]				
[0 0 0 0 0 5421 0 0 0 0]					[1 1 1 8 0 863 5 0 1 4]				
[0 0 0 0 0 0 5918 0 0 0]					[4 2 1 0 5 4 933 0 2 1]				
[0 0 0 0 0 0 0 6265 0 0]					[1 1 4 2 0 1 0 995 5 4]				
[0 0 0 0 0 0 0 0 5851 0]					[3 3 7 4 1 7 1 6 941 5]				
[0 0 0 0 0 0 0 0 0 5949]]					[1 0 0 9 9 3 0 6 6 982]]				
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	1.00	1.00	1.00	5923	0	0.99	0.98	0.98	992
1	1.00	1.00	1.00	6742	1	0.99	0.99	0.99	1136
2	1.00	1.00	1.00	5958	2	0.97	0.97	0.97	1036
3	1.00	1.00	1.00	6131	3	0.97	0.97	0.97	1008
4	1.00	1.00	1.00	5842	4	0.98	0.98	0.98	985
5	1.00	1.00	1.00	5421	5	0.97	0.98	0.97	884
6	1.00	1.00	1.00	5918	6	0.97	0.98	0.98	952
7	1.00	1.00	1.00	6265	7	0.97	0.98	0.98	1013
8	1.00	1.00	1.00	5851	8	0.97	0.96	0.96	978
9	1.00	1.00	1.00	5949	9	0.97	0.97	0.97	1016
accuracy			1.00	60000	accuracy			0.97	10000
macro avg	1.00	1.00	1.00	60000	macro avg	0.97	0.97	0.97	10000
weighted avg	1.00	1.00	1.00	60000	weighted avg	0.97	0.97	0.97	10000

- Then, I decided to use Randomized Search to find optimal hyperparameters, and to also compare these with ones chosen for the above models. I will use these parameters in model 2 to observe performance.

```
param_grid = {
    'hidden_layer_sizes':[32,64,96,128],
    'activation':['identity','logistic','tanh','relu'],
    'learning_rate':['constant', 'invscaling', 'adaptive'],
    'solver':['lbfgs','sgd','adam'],
    'batch_size':[100,250,500]
}
```

- Following parameters were picked:

```
[Parallel(n_jobs=1)]: Done 70 out of 70 | elapsed: 177.4min finished
best score 0.9776428571428573
best score {'solver': 'adam', 'learning_rate': 'constant', 'hidden_layer_sizes': 128, 'batch_size': 250, 'activation': 'logistic'}
```

```
Test data stats:
[[ 972  1  3  0  1  3  5  1  5  1]
 [ 0 1122  0  0  0  1  2  2  0  2]
 [ 1  4 1009  8  0  0  1 10  2  0]
 [ 1  0  5 984  1  8  1  2  4  4]
 [ 2  0  1  0 961  3  3  1  6 11]
 [ 0  1  0  5  0 867  5  1  4  2]
 [ 2  2  4  0  6  5 940  0  5  1]
 [ 1  1  5  5  0  0 1005  5  3]
 [ 1  4  5  3  2  3  1  2 941  5]
 [ 0  0  0  5 11  2  0  4  2 980]]
      precision    recall  f1-score   support

0         0.99      0.98      0.99         992
1         0.99      0.99      0.99        1129
2         0.98      0.97      0.98        1035
3         0.97      0.97      0.97        1010
4         0.98      0.97      0.98         988
5         0.97      0.98      0.98         885
6         0.98      0.97      0.98         965
7         0.98      0.98      0.98        1025
8         0.97      0.97      0.97         967
9         0.97      0.98      0.97        1004

 accuracy          0.98        10000
 macro avg          0.98        10000
 weighted avg       0.98        10000
```

- As you can see, after parameter search, model 2's accuracy slightly improved but this was at the cost of additional processing as it took almost 3 hours to find these parameters. A bigger problem or list of parameters to search through will take exponentially longer.

Model Comparison and Critical analysis about possible causes for inaccuracies

- **Scores:** Model 2 (scikit-learn) overall had a better accuracy than model 1. This could be due to various reasons such as underfitting model 1. Model 1 achieved the best scores with batch implemented, however training scores did not exceed 95% and there are still misclassifications present. Perhaps training model 1 more could result in scores closer to those seen with model 2.
- **Time:** In terms of training time, model 1 variants generally took longer to run than MLP even with batch implemented. This could be because the library has more complex logic implemented such as solvers and adaptive learning rates which could more efficiently train model 2.